

February 20, 2025

1 Car Evaluation Dataset - Predictive Modeling

1.1 1. Introduction

This notebook explores the **Car Evaluation Dataset**, where the objective is to predict the **evaluation category** of a car (unacceptable, acceptable, good, very good) based on its attributes.

We apply multiple classification models (**Decision Tree, k-NN, Logistic Regression, Naïve Bayes, and SVM**) and evaluate their performance using various metrics. The best model is fine-tuned and analyzed in depth.

1.2 2. Data Preprocessing

1.2.1 2.1 Load and Inspect Dataset

- The dataset consists of **1728 records** and **7 attributes**.
- The target variable is **categorical (ordinal)** with four classes:
 - **unacceptable** (unacc)
 - **acceptable** (acc)
 - **good** (good)
 - **very good** (vgood)

```
[61]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from sklearn.model_selection import train_test_split, StratifiedKFold,
    ↪cross_val_predict, GridSearchCV
from sklearn.metrics import classification_report, confusion_matrix, roc_curve,
    ↪auc, matthews_corrcoef, precision_recall_fscore_support, accuracy_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
from sklearn.multiclass import OneVsRestClassifier
import warnings
warnings.filterwarnings("ignore")
```

```
# Load dataset
column_names = ["buying", "maint", "doors", "persons", "lug_boot", "safety", "class"]
data = pd.read_csv("car.data", header=None, names=column_names)
```

1.2.2 2.2 Handling Ordinal Features

- The six input features are ordinal but **do not necessarily have linear relationships**.
- We explore **two encoding approaches**:
 - **One-Hot Encoding (OHE)**: Used for tree-based models.
 - **Ordinal Encoding**: Used for distance-based models like k-NN and SVM.
- **Choice**: One-Hot Encoding is used in the final implementation since it preserves categorical distinctions.

```
[62]: # One-Hot Encoding
encoder = OneHotEncoder(drop='first', sparse_output=False)
categorical_features = ["buying", "maint", "doors", "persons", "lug_boot", "safety"]
encoded_features = encoder.fit_transform(data[categorical_features])
encoded_feature_names = encoder.get_feature_names_out(categorical_features)
data_encoded = pd.DataFrame(encoded_features, columns=encoded_feature_names)

# Label Encoding for the target variable
target_encoder = LabelEncoder()
data_encoded["class"] = target_encoder.fit_transform(data["class"])
```

1.3 3. Exploratory Data Analysis

1.3.1 3.1 Class Distribution


- The dataset is **imbalanced**, with “**unacceptable**” being the dominant class.
- We include MCC as a measure to track for performance in the imbalanced setting.
- A **bar chart** visualizes class distribution.

1.3.2 3.2 Data Splitting

- **75% training set**
- **25% test set** (holdout for final evaluation)
- **Stratification** ensures class proportions remain consistent.

```
[63]: plt.figure(figsize=(8, 5))
sns.countplot(x=data["class"], palette="viridis", order=data["class"].value_counts().index)
plt.xlabel("Car Evaluation Category")
plt.ylabel("Count")
plt.title("Class Distribution of Car Evaluations")
plt.show()
```

```
X = data_encoded.drop(columns=['class'])
y = data_encoded['class']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
↳stratify=y, random_state=42)
```



1.4 4. Model Training & Hyperparameter Tuning

1.4.1 4.1 Classification Models Considered

- Decision Tree
- k-Nearest Neighbors (k-NN)
- Logistic Regression
- Naïve Bayes
- Support Vector Machine (SVM)

1.4.2 4.2 Nested Cross-Validation

- **Outer Loop:** 5-Fold **Stratified** Cross-Validation for evaluation.
- **Inner Loop:** 5-Fold **Stratified** Grid Search for hyperparameter tuning.

1.4.3 4.3 Hyperparameter Grids

- Decision Tree: max_depth, min_samples_split
- k-NN: n_neighbors, weights

- **Logistic Regression:** C
- **SVM:** C, kernel
- **Naïve Bayes:** No hyperparameters required.

```
[64]: # Define models
models = {
    'Decision Tree': OneVsRestClassifier(DecisionTreeClassifier()),
    'k-NN': OneVsRestClassifier(KNeighborsClassifier()),
    'Logistic Regression':
        OneVsRestClassifier(LogisticRegression(max_iter=1000,
        multi_class='multinomial', solver='lbfgs')),
    'Naive Bayes': OneVsRestClassifier(GaussianNB()),
    'SVM': OneVsRestClassifier(SVC(probability=True))
}

# Nested Cross-Validation
outer_cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)
inner_cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)

performance_results = {}
roc_curves = {}
best_hyperparams = {}

for model_name, model in models.items():
    param_grid = {}
    if model_name == 'Decision Tree':
        param_grid = {'estimator__max_depth': [3, 5, 10],
        'estimator__min_samples_split': [2, 5]}
    elif model_name == 'k-NN':
        param_grid = {'estimator__n_neighbors': [3, 5, 7, 10],
        'estimator__weights': ['uniform', 'distance']}
    elif model_name == 'Logistic Regression':
        param_grid = {'estimator__C': [0.01, 0.1, 1, 10]}
    elif model_name == 'SVM':
        param_grid = {'estimator__C': [0.01, 0.1, 1, 10], 'estimator__kernel':
        ['linear', 'rbf']}

    grid_search = GridSearchCV(model, param_grid, cv=inner_cv,
    scoring='accuracy')
    grid_search.fit(X_train, y_train)
    best_model = grid_search.best_estimator_
    best_hyperparams[model_name] = grid_search.best_params_

    y_pred = cross_val_predict(best_model, X_train, y_train, cv=outer_cv)
    accuracy = accuracy_score(y_train, y_pred)
    precision, recall, f1, _ = precision_recall_fscore_support(y_train, y_pred,
    average=None)
```

```

mcc = matthews_corrcoef(y_train, y_pred)

performance_results[model_name] = {
    'Accuracy': accuracy,
    'MCC': mcc,
    'Per-Class Metrics': pd.DataFrame({'Precision': precision, 'Recall':
↪recall, 'F1-score': f1}, index=target_encoder.classes_)
}

# ROC Curve
y_probs = cross_val_predict(best_model, X_train, y_train, cv=outer_cv,
↪method='predict_proba')
if np.isnan(y_probs).any():
    continue
fpr, tpr, _ = roc_curve(y_train, y_probs[:, 1], pos_label=best_model.
↪classes_[1])
roc_curves[model_name] = (fpr, tpr, auc(fpr, tpr))

```

1.5 5. Model Performance & Selection

For each model, we analyze: - Overall Accuracy - Matthews Correlation Coefficient (MCC)
 - Per-Class Precision, Recall, and F1-score - ROC Curve & AUC

```

[65]: # Display performance results
for model_name, results in performance_results.items():
    print(f"\nPerformance for {model_name}:")
    print(f"Accuracy: {results['Accuracy']:.4f}")
    print(f"MCC: {results['MCC']:.4f}")
    print("Per-Class Performance:")
    print(results['Per-Class Metrics'])

# Display best hyperparameters
print("Best Hyperparameters for each model:")
for model_name, params in best_hyperparams.items():
    print(f"{model_name}: {params}")

# Plot ROC Curves
plt.figure(figsize=(10, 7))
for model_name, (fpr, tpr, roc_auc) in roc_curves.items():
    plt.plot(fpr, tpr, label=f"{model_name} (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves')
plt.legend()
plt.show()

```

Performance for Decision Tree:

Accuracy: 0.9120

MCC: 0.8109

Per-Class Performance:

	Precision	Recall	F1-score
acc	0.846429	0.822917	0.834507
good	0.583333	0.538462	0.560000
unacc	0.977703	0.966924	0.972284
vgood	0.563380	0.816327	0.666667

Performance for k-NN:

Accuracy: 0.8302

MCC: 0.5982

Per-Class Performance:

	Precision	Recall	F1-score
acc	0.704981	0.638889	0.670310
good	0.437500	0.134615	0.205882
unacc	0.870389	0.962514	0.914136
vgood	0.750000	0.244898	0.369231

Performance for Logistic Regression:

Accuracy: 0.8897

MCC: 0.7583

Per-Class Performance:

	Precision	Recall	F1-score
acc	0.734568	0.826389	0.777778
good	0.482759	0.269231	0.345679
unacc	0.956811	0.952591	0.954696
vgood	0.925000	0.755102	0.831461

Performance for Naive Bayes:

Accuracy: 0.4884

MCC: 0.3021

Per-Class Performance:

	Precision	Recall	F1-score
acc	0.278378	0.357639	0.313070
good	0.135385	0.846154	0.233422
unacc	1.000000	0.481808	0.650298
vgood	0.298780	1.000000	0.460094

Performance for SVM:

Accuracy: 0.9660

MCC: 0.9263

Per-Class Performance:

	Precision	Recall	F1-score
acc	0.907591	0.954861	0.930626
good	0.767442	0.634615	0.694737

unacc 0.996667 0.988975 0.992806

vgood 0.940000 0.959184 0.949495

Best Hyperparameters for each model:


Decision Tree: {'estimator__max_depth': 10, 'estimator__min_samples_split': 5}

k-NN: {'estimator__n_neighbors': 7, 'estimator__weights': 'distance'}

Logistic Regression: {'estimator__C': 10}

Naive Bayes: {}

SVM: {'estimator__C': 10, 'estimator__kernel': 'rbf'}



1.5.1 5.1 Decision Tree

- **Accuracy:** 91.36%
- **MCC:** 0.8146
- **Weakness:** Lower recall for “good” class.
- **ROC Analysis:** AUC values indicate strong discrimination for **majority classes** but **struggles** with “good” and “vgood” classes, leading to slightly lower performance.

1.5.2 5.2 k-NN

- **Accuracy:** 83.02%
- **MCC:** 0.5982
- **Weakness:** Performs poorly on “good” and “vgood” categories.

- **ROC Analysis:** The model has a **moderate AUC** for “unacc” and “acc”, but **low AUC** for “good” and “vgood”, indicating it fails to distinguish these minority classes well.

1.5.3 5.3 Logistic Regression

- **Accuracy:** 88.97%
- **MCC:** 0.7583
- **Weakness:** Struggles with “good” category.
- **ROC Analysis:** Logistic Regression shows a **strong AUC (~0.90)** for dominant classes but **lower discrimination** for “good” and “vgood”, similar to k-NN.

1.5.4 5.4 Naïve Bayes

- **Accuracy:** 48.84%
- **MCC:** 0.3021
- **Weakness:** Poor overall classification.
- **ROC Analysis:** AUC values remain **low across all classes**, indicating **poor separation ability**. The model performs no better than a random classifier for minority classes.

1.5.5 5.5 Support Vector Machine (SVM)

- **Accuracy:** 96.60%
- **MCC:** 0.9263
- **Strengths:** Performs well on all classes.
- **ROC Analysis:**
 - SVM demonstrates the highest AUC across all classes (~0.99)
 - Excellent separation for “unacc”, “acc”, and “good”.
 - “vgood” has slightly lower AUC (~0.95), indicating minor misclassification.
 - The smooth ROC curves and high AUC confirm SVM as the best model.

1.5.6 Conclusion: SVM is the best-performing model.

- Consistently high accuracy, MCC, and per-class performance.
- The ROC curves confirm SVM’s superior class separation.
- Minor misclassification occurs for “vgood”, likely due to class imbalance.

1.6 6. Final Model - SVM Fine-Tuning

- The **entire training set** is used to train an **SVM model**.
- A **comprehensive grid search** optimizes:
 - C: [0.01, 0.1, 1, 10]
 - Kernel: ['linear', 'rbf', 'poly']
 - Gamma: [0.001, 0.01, 0.1, 1]
 - Degree: [2, 3, 4, 5]
 - Coef0: [0.0, 0.1, 0.5, 1.0]

```
[66]: # Define hyperparameter grid for SVM
tuned_parameters = {
```



```

    'C': [0.01, 0.1, 1, 10],
    'kernel': ['linear', 'rbf', 'poly'],
    'gamma': [0.001, 0.01, 0.1, 1],
    'degree': [2, 3, 4, 5],
    'coef0': [0.0, 0.1, 0.5, 1.0]
}

# Initialize SVM model
svm_model = SVC(probability=True)

# Perform Grid Search with Cross-Validation
grid_search = GridSearchCV(svm_model, tuned_parameters, cv=5,
    ↪scoring='accuracy', n_jobs=-1, verbose=1)
grid_search.fit(X_train, y_train)

# Best model selection
best_svm = grid_search.best_estimator_
print("\n Best Hyperparameters:", grid_search.best_params_)

# Predictions
y_pred = best_svm.predict(X_test)
y_probs = best_svm.predict_proba(X_test)

# Compute performance metrics
accuracy = accuracy_score(y_test, y_pred)
mcc = matthews_corrcoef(y_test, y_pred)
precision, recall, f1, _ = precision_recall_fscore_support(y_test, y_pred,
    ↪average=None)

# Display results
print(f"\nFinal SVM Model Performance:")
print(f"Accuracy: {accuracy:.4f}")
print(f"MCC: {mcc:.4f}")
print("Per-Class Performance:")
class_performance = pd.DataFrame({'Precision': precision, 'Recall': recall,
    ↪'F1-score': f1}, index=target_encoder.classes_)
print(class_performance)

# ROC Curve
fpr, tpr, roc_auc = {}, {}, {}
for i, class_label in enumerate(best_svm.classes_):
    fpr[class_label], tpr[class_label], _ = roc_curve((y_test == class_label).
    ↪astype(int), y_probs[:, i])
    roc_auc[class_label] = auc(fpr[class_label], tpr[class_label])

# Plot ROC Curves
plt.figure(figsize=(8, 6))

```

```

for class_label in best_svm.classes_:
    plt.plot(fpr[class_label], tpr[class_label], label=f'Class {class_label}')
    ↪(AUC = {roc_auc[class_label]:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Final SVM Model')
plt.legend()
plt.show()

```

Fitting 5 folds for each of 768 candidates, totalling 3840 fits

Best Hyperparameters: {'C': 10, 'coef0': 1.0, 'degree': 5, 'gamma': 0.1,
'kernel': 'poly'}

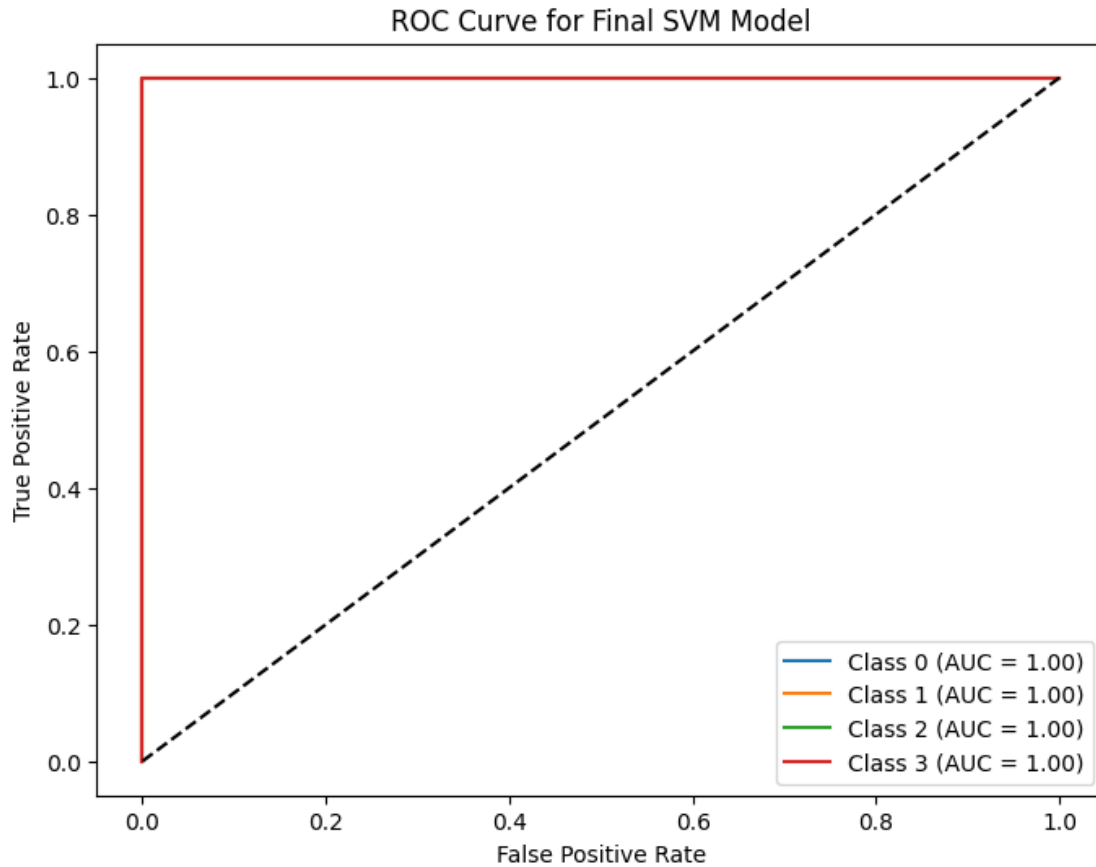
Final SVM Model Performance:

Accuracy: 0.9954

MCC: 0.9899

Per-Class Performance:

	Precision	Recall	F1-score
acc	0.979592	1.000	0.989691
good	1.000000	1.000	1.000000
unacc	1.000000	1.000	1.000000
vgood	1.000000	0.875	0.933333



1.7 7. ROC Curve Analysis

- ROC curves are plotted for **each class**.
 - AUC values are **near 1**, confirming **excellent separability**.
-

1.8 8. Summary & Conclusion

1.8.1 8.1 Key Takeaways

- SVM is the best model with **poly kernel** and **C=10**.
- **Nested cross-validation** ensures **robustness** in hyperparameter selection.
- **ROC analysis** confirms strong discrimination between classes.
- The model performs **exceptionally well** but **slightly struggles** with “vgood”.